

Concurrent and Distributed Programming (CA4006)

A Design Pattern Based Approach to Concurrency and Parallelization (Part 3)

2022/2023

Graham Healy

These course slides are partly adapted from the original course slides prepared by:
Dr Martin Crane, Dr Rob Brennan and Dr Takfarinas Saber

Introduction to OpenMP

- Stands for *Open Multi-Processing*, or *Open specifications for Multi-Processing*
- Represents collaboration between interested parties from h/w and s/w - industry, government and academia.
- An API to facilitate explicitly direct multi-threaded, shared memory parallelism.
- Supported in C, C++, and Fortran, and on most processor architectures and OS.
- Comprises a set of compiler directives, library routines, and environment variables affecting run-time behaviour.
- Introduce it here as complementary to and usable in conjunction with MPI (more later on this) to achieve speedup

Motivations to use OpenMP



- Provides a standard among a variety of shared memory architectures/platforms.
 - Currently at OpenMP Version 5.1 stable (as of Nov 2020)
 - More details at openmp.org/resources/
- Establishes a simple and limited set of directives for programming shared memory machines.
 - (like MPI) we can get quite good parallelism using 3 or 4 directives ...
- Unlike MPI:
 - Facilitates incremental parallelization of a serial program,
 - Does not require 'all or nothing' approach to parallelization,
 - MPI scales well but is non-trivial to implement for codes originally written for serial machines & not good for shared memory

What OpenMP is not

- OpenMP Is not:
 - Meant for distributed memory parallel systems (by itself)
 - Guaranteed to make the most efficient use of shared memory
- NB OpenMP will not:
 - Check for data dependencies, data conflicts, race conditions, or deadlocks
 - Check for code lines that cause program to be classified as non-conforming
 - Cover compiler-generated automatic parallelization and directives to the compiler to assist such parallelization

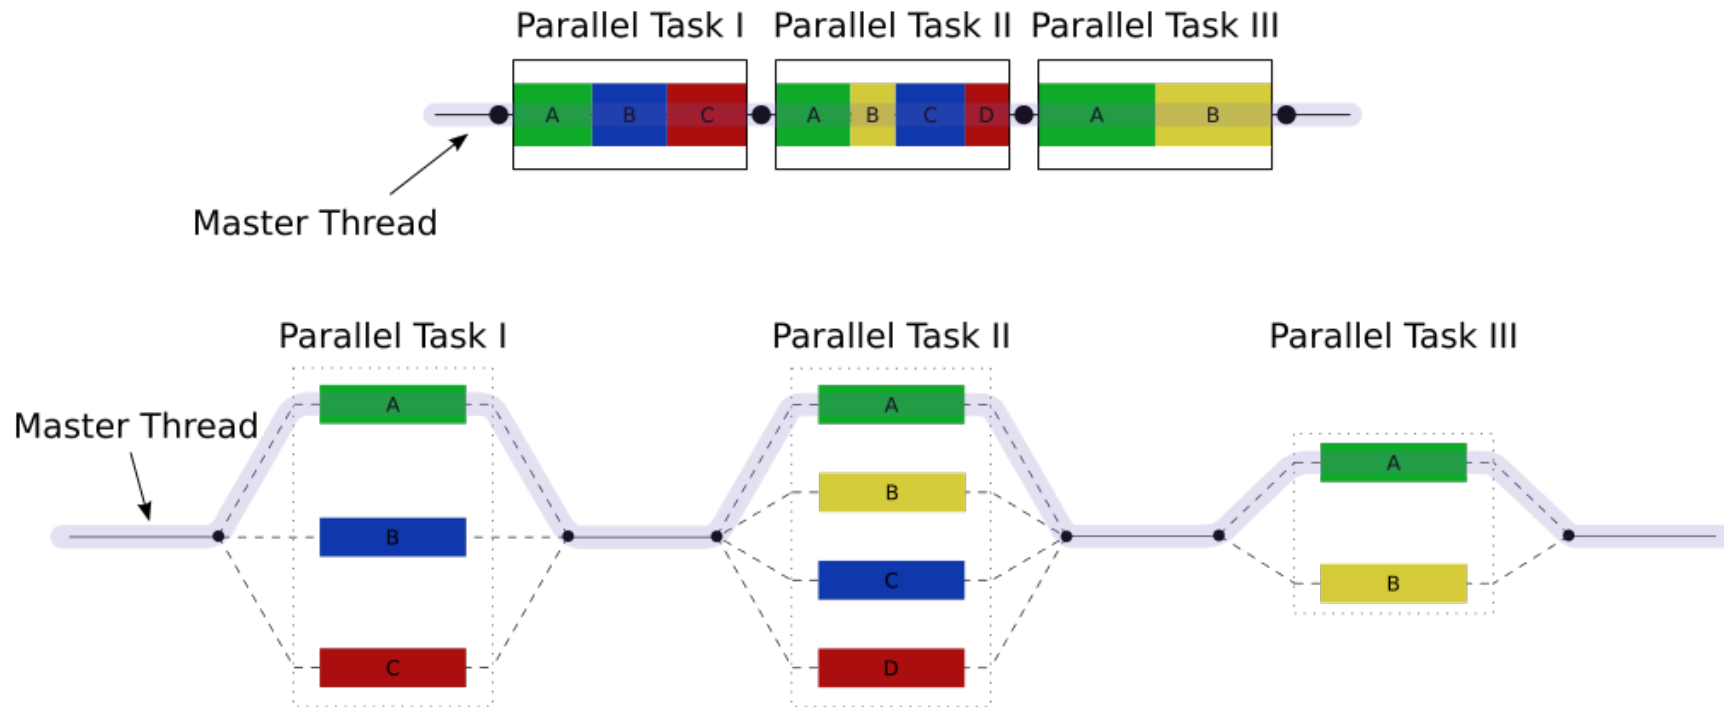
OpenMP Programming Model

- Shared Memory Model:
 - OpenMP is designed for multi-processor/core, shared memory machines.
 - The underlying architecture can be shared memory UMA or NUMA.



OpenMP Programming Model (/2)

The *Fork-Join* Model



"Fork join" by Wikipedia user A1 - w:en:File:Fork_join.svg. Licensed under CC BY 3.0 via Commons - https://commons.wikimedia.org/wiki/File:Fork_join.svg#/media/File:Fork_join.svg

Parallelism in OpenMP

- *Thread-Based Parallelism*

- OpenMP programs accomplish parallelism solely using *threads*.
- A *thread of execution* is smallest processing unit schedulable by OS.
 - Analogous conceptually to a subroutine that can be scheduled to run autonomously.
- These threads exist within resources of a single process, without which they cannot exist.
- Usually number of threads match the number of machine processors/cores.
 - However, the actual use of threads is up to the application.

Parallelism in OpenMP (/2)

- *Explicit Parallelism*

- OpenMP is an **explicit** (not automatic) programming model, offering the programmer full control over parallelization.
- Parallelization can be as simple as taking a serial program and inserting compiler directives....
- The general form of these are:

```
#pragma omp construct [clause [clause]...]
```

- Example of this:

```
#pragma omp parallel num_threads(4)
```

- Note about **#pragma**
 - These are special preprocessor instructions.
 - Typically added to system to allow behaviours that aren't part of the basic language specification.
 - Compilers that don't support the pragmas ignore them.

Example 1: My first OpenMP Code.

```
#include <stdio.h>
int main(void)
{
    #pragma omp parallel num_threads(2)
    printf("Hello, world.\n");
    return 0;
}
```

- *Thread Creation*

- `pragma omp parallel` used to fork additional threads (here 2) to carry out the work enclosed in the construct in parallel.
- The original thread is denoted as **master** thread with thread ID 0.
- `num_threads(2)` is one of a number of clauses that can be specified e.g. **private** variables, **shared** variables, **reduction** operation
- Simple Example: Display "Hello, world." using multiple threads.
- Complex: insert subroutines to set multiple levels of parallelism, locks and even nested locks.

Example 1: My first OpenMP Code (/2).

- *Thread Creation (/2)*

- When a thread reaches a `parallel` directive, creates a team of threads & becomes master of the team.
- From the start of this parallel region, code is duplicated and all threads execute that code.
- Implicit barrier at the end of a parallel section.
 - Only the master thread continues execution past this point.
- If any thread terminates in a parallel region, all threads in team stop
- If this happens, the work done up until that point is *undefined*.

Running this Example in OpenMP

- Use flag `-fopenmp` to compile using GCC:

```
$ gcc -fopenmp hello.c -o hello
```

- Outputs on a computer with 2 cores, and thus 2 threads:

```
Hello, world.
```

```
Hello, world.
```

- However, output may also be garbled due to race condition caused from the two threads sharing the standard output:

```
Hello, wHello, woorld.
```

```
rld.
```

- A helpful step by step example on how to run can be found at dartmouth.edu/~rc/classes/intro_openmp/compile_run.html

Example 2: More Complex OpenMP Code.

```
#include <stdio.h>
int main(int argc, char **argv) {
    int a[100];
    #pragma omp parallel for
    for (int i = 0; i < 100; i++)
        a[i] = 2 * i;
    return 0;
}
```

- Work-sharing constructs
 - `omp for/ omp do` for forking extra threads to do work enclosed in `parallel` (aka *loop* constructs).
 - This is equivalent to:

```
{    // stuff here                                }
#pragma omp parallel //nb omp twice

#pragma omp for (int i = 0; i < 100; i++)
    a[i] = 2 * i;
return 0;
```

Data Dependencies

- Data on one thread can be dependent on data on another one
- This can result in wrong answers
 - Thread 0 may require a variable that is calculated on thread 1
 - Answer depends on timing – When thread 0 does the calculation, has thread 1 calculated its value yet?
- Example – Fibonacci Sequence 0, 1, 1, 2, 3, 5, 8, 13, ... more bunnies!
- Parallelize on 2 threads
 - Thread 0 gets $i = 3$ to 51, Thread 1 gets $i = 52$ to 100
 - Look carefully at calculation for $i = 52$ on thread 1
 - What will be values of for $i - 1$ and $i - 2$?



```
A [1] = 0;  
A [2] = 1;  
for(i = 3; i <= 100; i++){  
    A [i] = A[i-1] + A[i-2];  
}
```

Data Dependencies (/2)

- *A Test for Dependency:*

- If serial loop is executed in reverse order, will it give same result?
- If so, it's ok
- You can test this on your serial code

- What about subprogram calls?

```
for(i = 0; i < 100; i++){  
    mycalc(i,x,y);  
}
```

Other Work Constructs in OpenMP

- **sections**

Used to assign consecutive but independent code blocks to different threads

- **single**

Specifying a code block that is executed by only one thread, a barrier is implied in the end

Uses first thread that encounters the construct.




- **master**

Similar to single, but code block is executed by **master** thread only – all others skip it

No barrier implied in the end.

OpenMP Work Constructs (Summary)

- **do / for** - shares loop iterations across team.
- Akin to "data parallelism"
- **sections** - breaks work into separate, discrete sections.
- Each executed by a thread.
- Can be used to implement a type of "functional parallelism".
- **single** - serializes a chunk of code



Example 3: **Sections** Construct.

```
void XAXIS(); void YAXIS(); void  
ZAXIS();  
void sect_example()  
{  
  #pragma omp parallel sections  
  {  
    #pragma omp section  
      XAXIS();  
    #pragma omp section  
      YAXIS();  
    #pragma omp section  
      ZAXIS();  
  }  
}
```

- Purpose
 - This **sections** directive is used to execute routines **XAXIS**, **YAXIS**, and **ZAXIS** concurrently

Example 4: **single** Construct.

```
#include <stdio.h>
void work1() {}
void work2() {}
void single_example() {
    #pragma omp parallel
    {
        #pragma omp single
        printf("Beginning work1.\n");
        work1();
        #pragma omp single
        printf("Finishing work1.\n");
        #pragma omp single nowait
        printf("Finished work1, starting work2.\n");
        work2();
    }
}
```

- Purpose
 - **single** directive specifies that the enclosed code is to be executed by only one thread in the team.
 - Useful dealing with sections of code that are not thread safe (such as I/O)
 - There is an implicit barrier at end of each except where a **nowait** clause is specified

Synchronisation Constructs in OpenMP

- **critical**
Specifies a critical section i.e. a region of code that must be executed by only one thread at a time.
- **atomic**
Commonly used to update counters and other simple variables that are accessed by multiple threads simultaneously.
- **barrier**
Synchronizes all threads in the team.
When reached, a thread waits there until all other threads have reached that barrier.
All then resume executing in parallel the code that follows the barrier.
- **master**
Strictly speaking, **master** is a synchronisation directive - master thread only and no barrier implied in the end.

Example 5: Data Scope Attributes

```
#include <stdio.h>
int a, b=0;
#pragma omp parallel for private(a) shared(b)
for(a=0; a<50; ++a)
{
    #pragma omp atomic // means that either happens or doesn't.
    b += a;           // one thread can't interrupt another here
}
```

- *Purpose*

- These attribute clauses specify data scoping/ sharing.
- As OpenMP based on shared memory programming model, most variables shared by default.
- Used with directives e.g. `Parallel`, `Do/ for`, `Sections` to control the scope of enclosed variables.
- `a` is explicitly specified `private` (each thread has own copy) and `b` is `shared` (each thread accesses same variable).

Example 6: A more Complex HelloWorld

```
#include <iostream>
using namespace std;
#include <omp.h>
int main(int argc, char *argv[])
{
    int th_id, nthreads;
    #pragma omp parallel private(th_id) shared(nthreads)
    {
        th_id = omp_get_thread_num(); // returns thread id
        #pragma omp critical // only one thread can access this at a time!
        {
            cout << "Hello World from thread " << th_id << '\n';
        }
        #pragma omp barrier // one thread waits for all others
        #pragma omp master // master thread access only!
        {
            nthreads = omp_get_num_threads(); // returns number of thread
            cout << "There are " << nthreads << " threads" << '\n';
        }
    }
    return 0;
}
```

- *Purpose*

- `Private`, `shared` declares that threads have their own copy of the variable or share a copy, respectively.

Reduction Clauses

- *Reduction*

- (Like MPI – you'll see later), OpenMP supports the Reduction operation.

```
int t;  
#pragma omp parallel reduction(+:t)  
{  
    t = omp_get_thread_num() + 1;  
    printf("local %d\n", t);  
}  
printf("reduction %d\n", t);
```

- Reduction Operators: **+** ***** **-** logical operators and **Min()**, **Max()**
- The operation makes the specified variable private to each thread.
- At the end of the computation it combines private results
- Very useful when combined with **for** as shown below see below:

```
sum = 0;  
#pragma omp parallel for reduction(+:sum)  
    for (i=0; i < 100; i++) {  
        sum += array[i];  
    }
```

Common Mistakes in OpenMP:

#1 Missing **Parallel** keyword

```
#pragma omp for //this is incorrect as parallel keyword omitted  
... // your code
```

- The code fragment will be successfully compiled, and the `#pragma omp for` directive will be simply ignored by the compiler.
- So only one thread executes the loop, and it could be tricky for a developer to uncover.
- The correct form should be:

```
#pragma omp parallel //this is correct  
{  
    #pragma omp for  
    ... //your code  
}
```

Common Mistakes in OpenMP:

#2 Missing **for** keyword

- **#pragma omp parallel**

- This directive may be applied to a single code line as well as to a code fragment. This may cause unexpected behaviour of the **for** loop:

```
#pragma omp parallel num_threads(2) // incorrect as for keyword omitted
for (int i = 0; i < 10; i++)
    myFunc();
```

- If the developer wanted to share the loop between two threads, they should use the **#pragma omp parallel for** directive.
- Here the loop would have been executed 10 times indeed.
- However, the code above will be executed once in every thread. As the result, the **myFunc();** function will be called 20 times.
- The correct version of the code is provided below:

```
#pragma omp parallel for num_threads(2) // now correct
for (int i = 0; i < 10; i++)
    myFunc();
```


Common Mistakes in OpenMP:

#3 Redundant Parallelization

- Applying the `#pragma omp parallel` directive to a large code fragment can lead to unexpected behaviour in cases like below:

```
#pragma omp parallel num_threads(2)
{
    ... // some lines of code
    #pragma omp parallel for
    for (int i = 0; i < 10; i++)
    {
        myFunc();
    }
}
```

- A naïve programmer wanting to share the loop execution between two threads placed the `parallel` keyword inside a parallel section.
- The result of execution is similar to previous example: the `myFunc` function will be called 20 times, not 10.
- The correct version of the code is the same as the above except for

```
#pragma omp parallel for
for (int i = 0; i < 10; i++)
```

Common Mistakes in OpenMP:

#4 Redefining the Number of Threads in a **Parallel** section

- By OpenMP 2.0 spec no. of threads cannot be redefined inside a **parallel** section without run-time errors and program termination:

```
#pragma omp parallel
{
  omp_set_num_threads(2); // incorrect to call this routine here
  #pragma omp for
  for (int i = 0; i < 10; i++)
    myFunc();
}
```

- A correct version of the code is:

```
#pragma omp parallel num_threads(2) // correct!
{
  #pragma omp for
  for (int i = 0; i < 10; i++)
    myFunc();
}
```

- Or:

```
omp_num_threads(2) // also
#pragma omp parallel // correct!
{
    #pragma omp for
    ... // more code
}
```